

Day 2 FAQ — Common Participant Questions

Quick answers to the questions that come up most. .NET AI with Claude Workshop.

Q1 “My Pro subscription works for Claude Code — why do I need an API key today?”

Pro / Max powers Claude.ai and Claude Code, but the **Anthropic API is billed separately** (pay-as-you-go). Day 2 calls the API directly, so it needs a **Console key**. New accounts get trial credit.

2 “Which NuGet package is the official one?”

Anthropic (v12.x). Community options exist (Anthropic.SDK, Claudia). Old **tryAGI.Anthropic v3.x** is a different lineage — migrate / upgrade.

3 “Why register AnthropicClient as a Singleton?”

It holds an **HttpClient** internally; Scoped / Transient leads to **socket exhaustion** — the same reason you use **IHttpClientFactory**.

4 “Do I need an OpenAI key for the RAG section?”

No — use the free / local path (Ollama + **nomic-embed-text**) behind the same **EmbeddingService** interface; set the Qdrant vector size to the model's dimensions.

5 “How is the C# MCP server different from yesterday's GitHub MCP server?”

Yesterday you **used** a server as a host; today you **build** one. Same architecture (host / client / server, tools / resources / prompts); you implement tools in C# and expose them over **Streamable HTTP** so any host — including Claude Code — can call them.

6 “How do I keep API costs under control?”

Prompt caching, batch processing, **Haiku** for mechanical / CI work, and per-request + per-user token budgets. See the Models & Cost Reference.

7 “How do I actually enforce our AI governance rules?”

Deterministic controls, not prompts — PreToolUse hooks, permissions, and managed settings (org-wide, non-overridable). Ties back to Day 1, Section 2.