

Models & Cost Reference

Free-first: default to free / local where possible; pay only where it's genuinely required.

1

Model lineup & pricing

per 1M tokens (USD)

MODEL	INPUT / OUTPUT	USE FOR
Claude Haiku 4.5	\$1 / \$5	High-volume mechanical work; CI code review
Claude Sonnet 4.6	\$3 / \$15	Default for app features
Claude Opus 4.8	\$5 / \$25	Hardest reasoning

Verify current rates before the event.

2

Cost-reduction levers

Prompt caching · Section 1

Up to ~90% off the cached input portion.

Batch processing

~50% cheaper for non-interactive workloads.

Model choice · Section 5

Haiku for CI reviews & mechanical tasks.

Token budgets · Section 6

Per-request MaxTokens cap + per-user daily budget.

3

Paid vs free / local map

CAPABILITY	NEEDS PAID?	FREE / LOCAL ALTERNATIVE
Claude API calls all sections	Yes — API key	Trial credit covers the workshop
RAG embeddings Section 3	OpenAI key by default	Ollama + nomic-embed-text (768-dim), same interface
Vector store Section 3	No	Qdrant in Docker (local)
Integration test infra Section 5	No	TestContainers (local Docker)
Claude Code Day 1	Pro min (\$20/mo)	— not available on Free

REMEMBER

A Claude.ai Pro / Max subscription does **NOT** grant API access — Day 2 needs a separate Anthropic Console key (pay-as-you-go).